Phase 2 Deliverables

Summary of Code-Base changes

Removing KV caching

llama/model.py:

- Commented out `kv_caching` flag in `ModelArgs`.
- Disabled all cache tensor pre-allocation and writes.
- Added non-cached fallback (`keys = xk`, `values = xv`).

llama/generation.py:

- Removed `kv_caching` argument from `Generation.generate()`.
- Simplified forward pass: always call `self(tokens[:, :cur_pos])`.
- Deleted cache-specific logic that updated `prev_pos`.

inference.py:

- Commented out old cache-enabled call; added new cache-free call.

Benchmark Testing

benchmark_inference.py:

- Updated call site to new `generate()` signature for testing purposes of this phase.

More detailed line-by-line changes can be viewed here:
https://github.com/uscmlsystems/ml-systems-final-project-Joshetaa/commit/09264895a68435e6d04ccc02746c7cc4d7a46a95

Comparison of model before and after KV cache removal

Before Cache removal

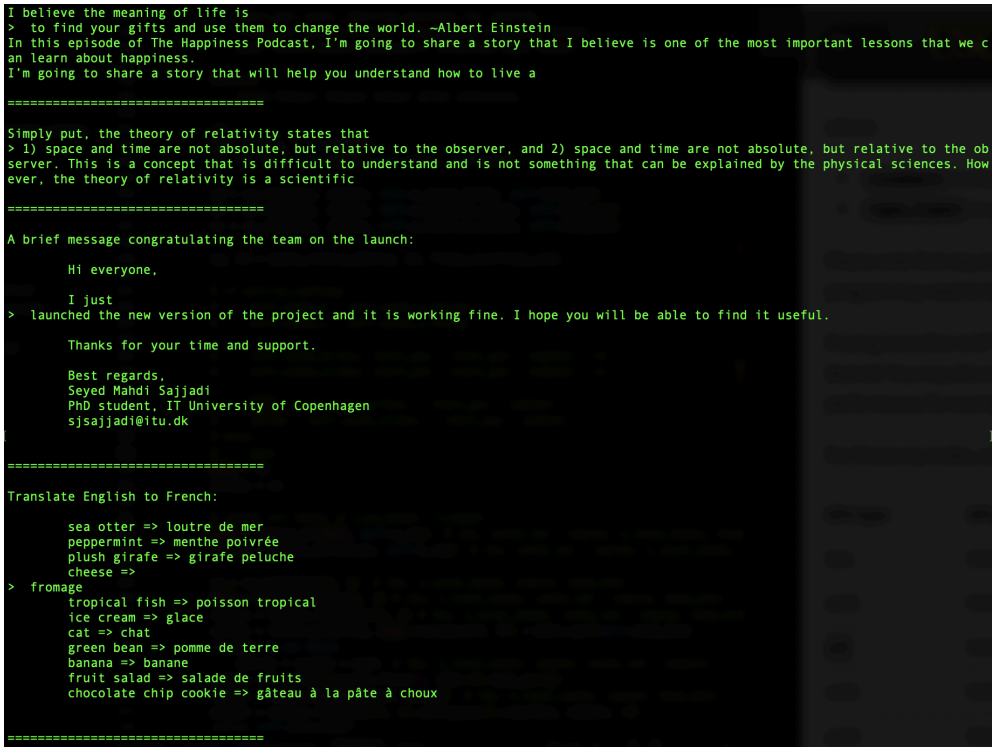

After Cache Removal

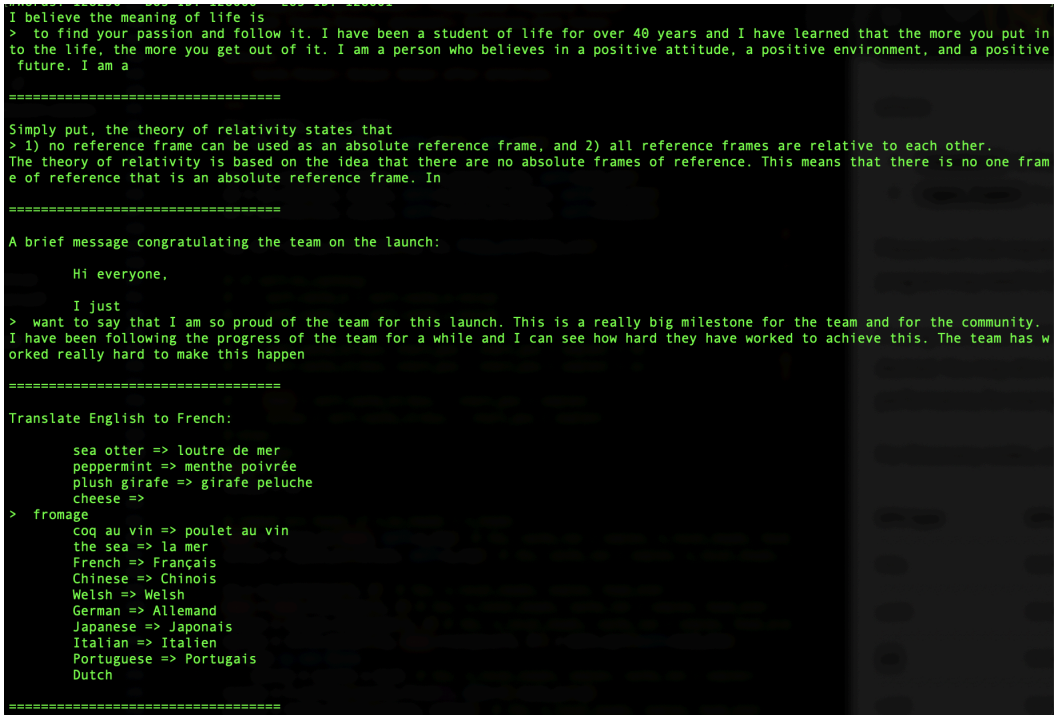

Qualitative Output Comparison (descriptive)

When KV-caching was enabled, the model's continuations tended toward longer, slightly more polished prose. For the prompt **"I believe the meaning of life is …"**, the cached run completed the sentence with an inspirational style that even quoted Einstein ("… to find your gifts and use them to change the world — Albert Einstein"). After removing the cache, the same prompt produced a shorter, more conversational answer ("… to find your passion and follow it. I have been a student of life for over 40 years…"). Both are on-topic, but the cache-free output feels less 'flowery' language.

A similar shift shows up in the physics prompt. With caching, the model began, "Simply put, the theory of relativity states that space and time are not absolute, but relative to the observer…," whereas the cache-free version reverted to the textbook framing about reference frames ("… no reference frame can be used as an absolute reference frame, and all frames are relative to each other…"). Content remains correct; only the phrasing differs. Although pre-cache removal, there seems to be redundant repetition.

For the **team-launch congratulatory email**, the cached model generated a full note complete with sign-off, name, and academic affiliation—clearly elaborated beyond the minimal ask. The cache-free model kept things shorter and more generic, omitting the formal signature block. Again, semantics align, but verbosity is reduced without the cache.

Finally, in the **English→French lexicon prompt** ("cheese ⇒ ?"), both versions correctly answered *fromage* and then free-associated additional word pairs. The cached run listed food items like *tropical fish* and *ice cream*, while the cache-free run chose examples such as *coq au vin* and language names (Welsh, German). The differences are arbitrary expansions rather than errors, suggesting here that disabling KV-caching affects stylistic sampling choices more than factual accuracy.

Benchmarking

input len=256, output len=32		batch size=1	batch size=8	batch size=16
with KV cache	Peak Mem	3071.57 MB	4495.47 MB	6133.80 MB
	Runtime	0.37 seconds	0.52 seconds	0.63 seconds
without KV cache	Peak Mem	3230.12 MB	5755.00 MB	8641.54 MB
	Runtime	0.41 seconds	1.80 seconds	4.25 seconds

1. Memory behaviour

At every batch size the cache-enabled run consumes **less** memory than the cache-free run.

- **Batch 1:** 3.07 GB with cache vs 3.23 GB without (≈ +5 %).
- **Batch 8:** 4.50 GB with cache vs 5.76 GB without (≈ +28 %).
- **Batch 16:** 6.13 GB with cache vs 8.64 GB without (≈ +41 %).

The gap widens as the batch size grows because, without caching, each forward step must materialise fresh key/value projections for *all* decoder layers, whereas the cached variant stores them once and simply appends new tokens. Consequently, memory overhead scales almost linearly with sequence-length × batch-size in the cache-free setting.

2. Runtime behaviour

Caching also delivers consistent speed-ups, and the advantage becomes bigger for larger batches:

- **Batch 1:** 0.37 s (cached) vs 0.41 s (cache-free) — a modest 11 % slowdown.
- **Batch 8:** 0.52 s vs 1.80 s — roughly **3.5 ×** slower without the cache.
- **Batch 16:** 0.63 s vs 4.25 s — roughly **6.7 ×** slower.

Because the cache-free path recomputes attention for the entire prefix at every decoding step, its cost grows with both *sequence length* and *batch size*. In contrast, the cached version reuses previously computed key/value tensors and only attends over the newly generated token, hence the nearly flat runtime curve.

My Take-aways

- **When VRAM is plentiful:** keep KV caching turned **on**—it both lowers peak memory and accelerates generation, with benefits that compound for larger batches or longer outputs.
- **When VRAM is extremely limited or on CPU:** disabling the cache can still work (no correctness loss), but expect noticeably higher latency; you are trading compute for model-state simplicity.
- The latency penalty grows super-linearly with batch size, so cache-free inference is best reserved for single-request, low-throughput scenarios.